

```

# -*- coding: utf-8 -*-
"""
CIFake Image Classification - Restructured Version

Maintains identical functionality to original script but uses modular functions.
"""

import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import kagglehub
import tensorflow as tf
from tensorflow.keras import layers, models, callbacks
from tensorflow.keras.preprocessing import image_dataset_from_directory
from sklearn.metrics import confusion_matrix

# Configuration constants
CONFIG = {
    "dataset_slug": "birdy654/cifake-real-and-ai-generated-synthetic-images",
    "image_size": (32, 32),
    "batch_size": 32,
    "validation_split": 0.2,
    "random_seed": 123,
    "autotune": tf.data.AUTOTUNE,
    "training_epochs": 10,
    "early_stopping_patience": 5,
    "model_save_path": "best_model.keras"
}

def handle_dataset_download():
    """Download and validate dataset structure"""
    try:
        base_path = kagglehub.dataset_download(CONFIG["dataset_slug"])
        train_path = os.path.join(base_path, 'train')
        test_path = os.path.join(base_path, 'test')

        if not all(os.path.exists(p) for p in [train_path, test_path]):
            raise FileNotFoundError("Dataset directories missing")

        return train_path, test_path

    except Exception as e:
        print(f"Dataset error: {str(e)}")
        exit()

def create_datasets(train_dir, test_dir):
    """Create optimized TensorFlow datasets"""
    try:
        train_ds = image_dataset_from_directory(
            train_dir,
            validation_split=CONFIG["validation_split"],
            subset="training",
            seed=CONFIG["random_seed"],
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"],
            label_mode='binary'
        )

        val_ds = image_dataset_from_directory(
            train_dir,
            validation_split=CONFIG["validation_split"],
            subset="validation",
            seed=CONFIG["random_seed"],
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"]
        )

        test_ds = image_dataset_from_directory(
            test_dir,
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"],
            shuffle=False
        )

    return (

```

```

        train_ds.prefetch(CONFIG["autotune"]),
        val_ds.prefetch(CONFIG["autotune"]),
        test_ds.prefetch(CONFIG["autotune"]),
        train_ds.class_names
    )

except Exception as e:
    print(f"Data loading error: {str(e)}")
    exit()

def build_cnn_model():
    """Construct the CNN architecture"""
    input_shape = CONFIG["image_size"] + (3,)

    model = models.Sequential([
        layers.Input(shape=input_shape),
        layers.Rescaling(1./255),
        layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2)),
        layers.Flatten(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(1, activation='sigmoid')
    ])

    model.compile(
        optimizer='adam',
        loss='binary_crossentropy',
        metrics=['accuracy']
    )
    return model

def train_model(model, train_data, val_data):
    """Handle model training with callbacks"""
    # Changed variable name to avoid conflict with callbacks module
    callback_list = [
        callbacks.EarlyStopping(
            monitor='val_loss',
            patience=CONFIG["early_stopping_patience"],
            restore_best_weights=True,
            verbose=1
        ),
        callbacks.ModelCheckpoint(
            CONFIG["model_save_path"],
            monitor='val_accuracy',
            save_best_only=True,
            verbose=1
        )
    ]

    return model.fit(
        train_data,
        validation_data=val_data,
        epochs=CONFIG["training_epochs"],
        callbacks=callback_list # Changed variable name here
    )

def visualize_results(history, class_names, test_data, model):
    """Generate performance visualizations"""
    # Training history plots
    history_df = pd.DataFrame(history.history)
    fig, ax = plt.subplots(1, 2, figsize=(14, 5))

    ax[0].plot(history_df['accuracy'], label='Train')
    ax[0].plot(history_df['val_accuracy'], label='Validation')
    ax[0].set_title('Accuracy Metrics')
    ax[0].set_ylabel('Accuracy')
    ax[0].legend()

    ax[1].plot(history_df['loss'], label='Train')
    ax[1].plot(history_df['val_loss'], label='Validation')
    ax[1].set_title('Loss Metrics')
    ax[1].set_ylabel('Loss')
    ax[1].legend()

    plt.tight_layout()

```

```
plt.show()

# Confusion matrix
y_pred = (model.predict(test_data) > 0.5).astype(int).flatten()
y_true = np.concatenate([y.numpy() for _, y in test_data], axis=0)

plt.figure(figsize=(8, 6))
sns.heatmap(
    confusion_matrix(y_true, y_pred),
    annot=True,
    fmt='d',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

def main():
    """Main execution flow"""
    # Dataset setup
    global test_ds # Define test_ds as a global variable
    train_path, test_path = handle_dataset_download()
    train_ds, val_ds, test_ds, classes = create_datasets(train_path, test_path)


    # Model lifecycle
    model = build_cnn_model()
    model.summary()

    history = train_model(model, train_ds, val_ds)

    # Final evaluation
    test_results = model.evaluate(test_ds, verbose=0)
    print(f"\nTest Accuracy: {test_results[1]:.2%}")
    print(f"Test Loss: {test_results[0]:.4f}")

    # Visualizations
    visualize_results(history, classes, test_ds, model)

if __name__ == "__main__":
    main()
```

Warning: Looks like you're using an outdated `kagglehub` version (installed: 0.3.10), please consider upgrading to the latest version. 

Found 100000 files belonging to 2 classes.
 Using 80000 files for training.
 Found 100000 files belonging to 2 classes.
 Using 20000 files for validation.
 Found 20000 files belonging to 2 classes.
 Model: "sequential_3"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 32, 32, 3)	0
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_1 (Flatten)	(None, 4096)	0
dense_6 (Dense)	(None, 128)	524,416
dropout_3 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 1)	129

Total params: 543,937 (2.07 MB)

Trainable params: 543,937 (2.07 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.8168 - loss: 0.3951

Epoch 1: val_accuracy improved from -inf to 0.87445, saving model to best_model.keras

2500/2500 ————— 168s 66ms/step - accuracy: 0.8168 - loss: 0.3951 - val_accuracy: 0.8745 - val_loss: 0.304

Epoch 2/10

2500/2500 ————— 0s 61ms/step - accuracy: 0.9125 - loss: 0.2185

Epoch 2: val_accuracy improved from 0.87445 to 0.91375, saving model to best_model.keras

2500/2500 ————— 210s 69ms/step - accuracy: 0.9125 - loss: 0.2185 - val_accuracy: 0.9137 - val_loss: 0.224

Epoch 3/10

2500/2500 ————— 0s 60ms/step - accuracy: 0.9297 - loss: 0.1829

Epoch 3: val_accuracy improved from 0.91375 to 0.93515, saving model to best_model.keras

2500/2500 ————— 172s 69ms/step - accuracy: 0.9297 - loss: 0.1829 - val_accuracy: 0.9352 - val_loss: 0.168

Epoch 4/10

2499/2500 ————— 0s 61ms/step - accuracy: 0.9386 - loss: 0.1618

Epoch 4: val_accuracy improved from 0.93515 to 0.94010, saving model to best_model.keras

2500/2500 ————— 173s 69ms/step - accuracy: 0.9386 - loss: 0.1618 - val_accuracy: 0.9401 - val_loss: 0.162

Epoch 5/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.9447 - loss: 0.1455

Epoch 5: val_accuracy did not improve from 0.94010

2500/2500 ————— 195s 67ms/step - accuracy: 0.9447 - loss: 0.1455 - val_accuracy: 0.9366 - val_loss: 0.172

Epoch 6/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.9487 - loss: 0.1326

Epoch 6: val_accuracy improved from 0.94010 to 0.94035, saving model to best_model.keras

2500/2500 ————— 174s 70ms/step - accuracy: 0.9487 - loss: 0.1326 - val_accuracy: 0.9403 - val_loss: 0.169

Epoch 7/10

2500/2500 ————— 0s 61ms/step - accuracy: 0.9540 - loss: 0.1197

Epoch 7: val_accuracy improved from 0.94035 to 0.94380, saving model to best_model.keras

2500/2500 ————— 201s 70ms/step - accuracy: 0.9540 - loss: 0.1197 - val_accuracy: 0.9438 - val_loss: 0.154

Epoch 8/10

2500/2500 ————— 0s 61ms/step - accuracy: 0.9594 - loss: 0.1077

Epoch 8: val_accuracy did not improve from 0.94380

2500/2500 ————— 200s 69ms/step - accuracy: 0.9594 - loss: 0.1077 - val_accuracy: 0.9410 - val_loss: 0.163

Epoch 9/10

2499/2500 ————— 0s 62ms/step - accuracy: 0.9621 - loss: 0.0997

Epoch 9: val_accuracy improved from 0.94380 to 0.94435, saving model to best_model.keras

2500/2500 ————— 196s 67ms/step - accuracy: 0.9621 - loss: 0.0997 - val_accuracy: 0.9444 - val_loss: 0.165

Epoch 10/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.9646 - loss: 0.0916


Epoch 10: val_accuracy did not improve from 0.94435


2500/2500 ————— 200s 66ms/step - accuracy: 0.9646 - loss: 0.0916 - val_accuracy: 0.9441 - val_loss: 0.176

Restoring model weights from the end of the best epoch: 7.


Test Accuracy: 94.46%

Test Loss: 0.1542





Confusion Matrix



Start coding or [generate](#) with AI.

```
# -*- coding: utf-8 -*-
"""
```

CIFake Image Classification - Restructured Version

Maintains identical functionality to original script but uses modular functions.
"""

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import kagglehub
import tensorflow as tf
from tensorflow.keras import layers, models, callbacks
from tensorflow.keras.preprocessing import image_dataset_from_directory
from sklearn.metrics import confusion_matrix
```

Configuration constants

```
CONFIG = {
    "dataset_slug": "birdy654/cifake-real-and-ai-generated-synthetic-images",
    "image_size": (32, 32),
    "batch_size": 32,
    "validation_split": 0.2,
    "random_seed": 123,
    "autotune": tf.data.AUTOTUNE,
    "training_epochs": 10,
    "early_stopping_patience": 5,
    "model_save_path": "best_model.keras"
}
```

```
def handle_dataset_download():
    """Download and validate dataset structure"""
    try:
        base_path = kagglehub.dataset_download(CONFIG["dataset_slug"])
        train_path = os.path.join(base_path, 'train')
        test_path = os.path.join(base_path, 'test')

        if not all(os.path.exists(p) for p in [train_path, test_path]):
            raise FileNotFoundError("Dataset directories missing")

        return train_path, test_path

    except Exception as e:
        print(f"Dataset error: {str(e)}")
        exit()
```

```
def create_datasets(train_dir, test_dir):
    """Create optimized TensorFlow datasets"""
    try:
        train_ds = image_dataset_from_directory(
            train_dir,
            validation_split=CONFIG["validation_split"],
            subset="training",
            seed=CONFIG["random_seed"],
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"],
            label_mode='binary'
        )

        val_ds = image_dataset_from_directory(
            train_dir,
            validation_split=CONFIG["validation_split"],
            subset="validation",
            seed=CONFIG["random_seed"],
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"]
        )

        test_ds = image_dataset_from_directory(
            test_dir,
            image_size=CONFIG["image_size"],
            batch_size=CONFIG["batch_size"],
            shuffle=False
```

```

    )

    return (
        train_ds.prefetch(CONFIG["autotune"]),
        val_ds.prefetch(CONFIG["autotune"]),
        test_ds.prefetch(CONFIG["autotune"]),
        train_ds.class_names
    )

except Exception as e:
    print(f"Data loading error: {str(e)}")
    exit()

def build_cnn_model():
    """Construct the CNN architecture"""
    input_shape = CONFIG["image_size"] + (3,)

    model = models.Sequential([
        layers.Input(shape=input_shape),
        layers.Rescaling(1./255),
        layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2)),
        layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
        layers.MaxPooling2D((2, 2)),
        layers.Flatten(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(1, activation='sigmoid')
    ])

    model.compile(
        optimizer='adam',
        loss='binary_crossentropy',
        metrics=['accuracy']
    )
    return model

def train_model(model, train_data, val_data):
    """Handle model training with callbacks"""
    # Changed variable name to avoid conflict with callbacks module
    callback_list = [
        callbacks.EarlyStopping(
            monitor='val_loss',
            patience=CONFIG["early_stopping_patience"],
            restore_best_weights=True,
            verbose=1
        ),
        callbacks.ModelCheckpoint(
            CONFIG["model_save_path"],
            monitor='val_accuracy',
            save_best_only=True,
            verbose=1
        )
    ]

    return model.fit(
        train_data,
        validation_data=val_data,
        epochs=CONFIG["training_epochs"],
        callbacks=callback_list # Changed variable name here
    )

def visualize_results(history, class_names, test_data, model):
    """Generate performance visualizations"""
    # Training history plots
    history_df = pd.DataFrame(history.history)
    fig, ax = plt.subplots(1, 2, figsize=(14, 5))

    ax[0].plot(history_df['accuracy'], label='Train')
    ax[0].plot(history_df['val_accuracy'], label='Validation')
    ax[0].set_title('Accuracy Metrics')
    ax[0].set_ylabel('Accuracy')
    ax[0].legend()

    ax[1].plot(history_df['loss'], label='Train')
    ax[1].plot(history_df['val_loss'], label='Validation')
    ax[1].set_title('Loss Metrics')
    ax[1].set_ylabel('Loss')

```

```

ax[1].legend()

plt.tight_layout()
plt.show()

# Confusion matrix
y_pred = (model.predict(test_data) > 0.5).astype(int).flatten()
y_true = np.concatenate([y.numpy() for _, y in test_data], axis=0)

plt.figure(figsize=(8, 6))
sns.heatmap(
    confusion_matrix(y_true, y_pred),
    annot=True,
    fmt='d',
    xticklabels=class_names,
    yticklabels=class_names
)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

def generate_adversarial_examples(model, dataset, epsilon=0.01):
    """Generate adversarially perturbed images using FGSM"""
    adv_images = []
    adv_labels = []
    loss_object = tf.keras.losses.BinaryCrossentropy()

    for images, labels in dataset:
        images = tf.cast(images, tf.float32) / 255.0
        with tf.GradientTape() as tape:
            tape.watch(images)
            predictions = model(images)
            loss = loss_object(labels, predictions)

            gradients = tape.gradient(loss, images)
            perturbed_images = images + epsilon * tf.sign(gradients)
            perturbed_images = tf.clip_by_value(perturbed_images, 0, 1)

            adv_images.append(perturbed_images)
            adv_labels.append(labels)

    return tf.data.Dataset.from_tensor_slices((tf.concat(adv_images, axis=0), tf.concat(adv_labels, axis=0))).batch(CONFIG["batch

def main():
    """Main execution flow"""
    # Dataset setup
    global test_ds # Define test_ds as a global variable
    train_path, test_path = handle_dataset_download()
    train_ds, val_ds, test_ds, classes = create_datasets(train_path, test_path)

    # Model lifecycle
    model = build_cnn_model()
    model.summary()

    history = train_model(model, train_ds, val_ds)

    # Generate adversarial dataset
    adv_train_ds = generate_adversarial_examples(model, train_ds)


    # Train model again with adversarial dataset
    print("\nTraining model with adversarial examples...")
    history_adv = train_model(model, adv_train_ds, val_ds)

    # Final evaluation
    test_results = model.evaluate(test_ds, verbose=0)
    print(f"\nTest Accuracy: {test_results[1]:.2%}")
    print(f"Test Loss: {test_results[0]:.4f}")

    # Visualizations
    visualize_results(history, classes, test_ds, model)

if __name__ == "__main__":
    main()

```


Warning: Looks like you're using an outdated `kagglehub` version (installed: 0.3.10), please consider upgrading to the latest version. 

Found 100000 files belonging to 2 classes.
 Using 80000 files for training.
 Found 100000 files belonging to 2 classes.
 Using 20000 files for validation.
 Found 20000 files belonging to 2 classes.
 Model: "sequential_7"

Layer (type)	Output Shape	Param #
rescaling_5 (Rescaling)	(None, 32, 32, 3)	0
conv2d_6 (Conv2D)	(None, 32, 32, 32)	896
max_pooling2d_6 (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_7 (Conv2D)	(None, 16, 16, 64)	18,496
max_pooling2d_7 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten_3 (Flatten)	(None, 4096)	0
dense_14 (Dense)	(None, 128)	524,416
dropout_7 (Dropout)	(None, 128)	0
dense_15 (Dense)	(None, 1)	129

Total params: 543,937 (2.07 MB)

Trainable params: 543,937 (2.07 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.8221 - loss: 0.3869

Epoch 1: val_accuracy improved from -inf to 0.89675, saving model to best_model.keras

2500/2500 ————— 178s 70ms/step - accuracy: 0.8221 - loss: 0.3869 - val_accuracy: 0.8967 - val_loss: 0.254

Epoch 2/10

2500/2500 ————— 0s 63ms/step - accuracy: 0.9149 - loss: 0.2153

Epoch 2: val_accuracy improved from 0.89675 to 0.92530, saving model to best_model.keras

2500/2500 ————— 194s 67ms/step - accuracy: 0.9149 - loss: 0.2153 - val_accuracy: 0.9253 - val_loss: 0.194

Epoch 3/10

2499/2500 ————— 0s 62ms/step - accuracy: 0.9286 - loss: 0.1839

Epoch 3: val_accuracy improved from 0.92530 to 0.93030, saving model to best_model.keras

2500/2500 ————— 176s 71ms/step - accuracy: 0.9286 - loss: 0.1839 - val_accuracy: 0.9303 - val_loss: 0.181

Epoch 4/10

2500/2500 ————— 0s 64ms/step - accuracy: 0.9360 - loss: 0.1649

Epoch 4: val_accuracy improved from 0.93030 to 0.93850, saving model to best_model.keras

2500/2500 ————— 196s 68ms/step - accuracy: 0.9360 - loss: 0.1649 - val_accuracy: 0.9385 - val_loss: 0.161

Epoch 5/10

2499/2500 ————— 0s 63ms/step - accuracy: 0.9437 - loss: 0.1488

Epoch 5: val_accuracy did not improve from 0.93850

2500/2500 ————— 170s 68ms/step - accuracy: 0.9438 - loss: 0.1488 - val_accuracy: 0.9330 - val_loss: 0.189

Epoch 6/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.9488 - loss: 0.1341

Epoch 6: val_accuracy improved from 0.93850 to 0.94155, saving model to best_model.keras

2500/2500 ————— 199s 67ms/step - accuracy: 0.9488 - loss: 0.1341 - val_accuracy: 0.9416 - val_loss: 0.161

Epoch 7/10

2500/2500 ————— 0s 62ms/step - accuracy: 0.9532 - loss: 0.1206

Epoch 7: val_accuracy improved from 0.94155 to 0.94475, saving model to best_model.keras

2500/2500 ————— 177s 71ms/step - accuracy: 0.9532 - loss: 0.1206 - val_accuracy: 0.9448 - val_loss: 0.154

Epoch 8/10

2500/2500 ————— 0s 64ms/step - accuracy: 0.9584 - loss: 0.1106

Epoch 8: val_accuracy did not improve from 0.94475

2500/2500 ————— 205s 72ms/step - accuracy: 0.9584 - loss: 0.1106 - val_accuracy: 0.9366 - val_loss: 0.169

Epoch 9/10

2500/2500 ————— 0s 63ms/step - accuracy: 0.9606 - loss: 0.1043

Epoch 9: val_accuracy improved from 0.94475 to 0.94585, saving model to best_model.keras

2500/2500 ————— 201s 72ms/step - accuracy: 0.9606 - loss: 0.1043 - val_accuracy: 0.9459 - val_loss: 0.157

Epoch 10/10

2500/2500 ————— 0s 63ms/step - accuracy: 0.9640 - loss: 0.0956

Epoch 10: val_accuracy did not improve from 0.94585

2500/2500 ————— 168s 67ms/step - accuracy: 0.9640 - loss: 0.0956 - val_accuracy: 0.9452 - val_loss: 0.166

Restoring model weights from the end of the best epoch: 9.

Training model with adversarial examples...

Epoch 1/10

2500/2500 ————— 0s 59ms/step - accuracy: 0.4930 - loss: 0.9245

Epoch 1: val_accuracy improved from -inf to 0.89780, saving model to best_model.keras

2500/2500 ————— 160s 64ms/step - accuracy: 0.4930 - loss: 0.9244 - val_accuracy: 0.8978 - val_loss: 0.266

Epoch 2/10

2500/2500 ————— 0s 58ms/step - accuracy: 0.4971 - loss: 0.6932

Epoch 2: val_accuracy did not improve from 0.89780

2500/2500 ————— 199s 62ms/step - accuracy: 0.4971 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.266

Epoch 3/10

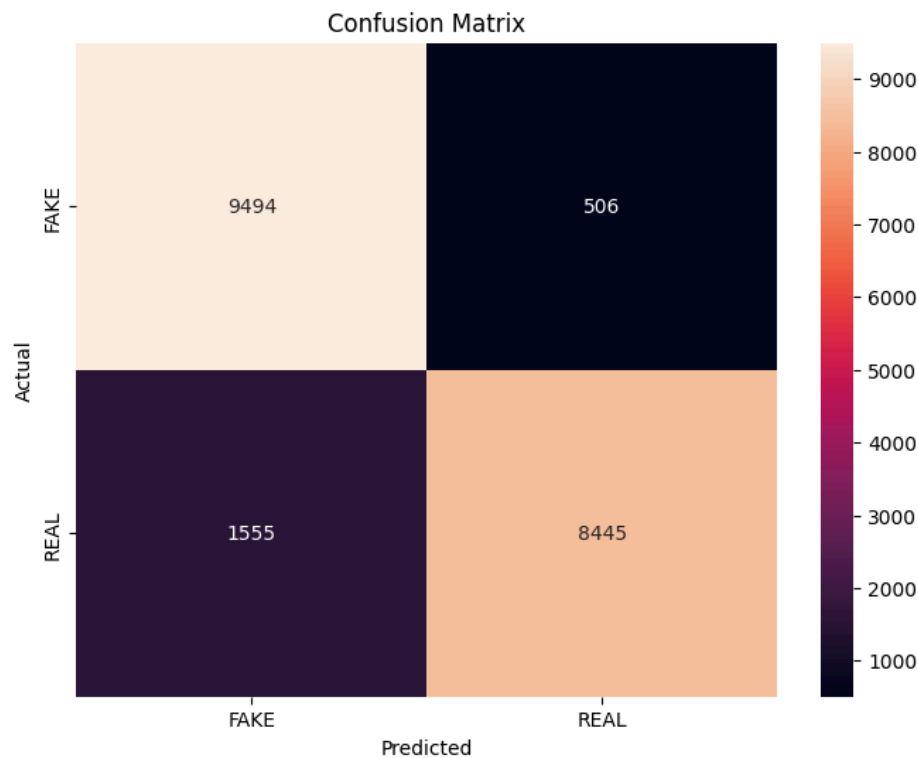

2500/2500 ————— 0s 58ms/step - accuracy: 0.4966 - loss: 0.6932

Epoch 3: val_accuracy did not improve from 0.89780

2500/2500 ————— 155s 66ms/step - accuracy: 0.4966 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.266

```
2500/2500 ————— 103s 60ms/step - accuracy: 0.4966 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.2546
Epoch 4/10
2500/2500 ————— 0s 58ms/step - accuracy: 0.4966 - loss: 0.6932
Epoch 4: val_accuracy did not improve from 0.89780
2500/2500 ————— 155s 62ms/step - accuracy: 0.4966 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.2546
Epoch 5/10
2500/2500 ————— 0s 59ms/step - accuracy: 0.4966 - loss: 0.6932
Epoch 5: val_accuracy did not improve from 0.89780
2500/2500 ————— 167s 67ms/step - accuracy: 0.4966 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.2546
Epoch 6/10
2500/2500 ————— 0s 58ms/step - accuracy: 0.4966 - loss: 0.6932
Epoch 6: val_accuracy did not improve from 0.89780
2500/2500 ————— 157s 63ms/step - accuracy: 0.4966 - loss: 0.6932 - val_accuracy: 0.8978 - val_loss: 0.2546
Epoch 6: early stopping
Restoring model weights from the end of the best epoch: 1.
```

Test Accuracy: 89.70%
Test Loss: 0.2546



```
import zipfile
import os

zip_path = "/content/Test datasets.zip" # Replace with your actual zip file path
extract_path = "/content" # Change this to your desired output folder

with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_path)

print("Unzipping Done!")
```

➞ Unzipping Done!

```
import os
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import pandas as pd # Added for CSV generation
from tensorflow.keras.preprocessing import image

# Load the trained model
model = tf.keras.models.load_model(CONFIG["model_save_path"])

# Paths to test datasets
testdataset1_path = "/content/Test datasets/Test_dataset_1"

# Function to load, preprocess, predict, and save results
def predict_images_from_folder(folder_path, output_csv="predictions_test1.csv"):
    """Loads images, preprocesses them, predicts labels, and saves results to CSV."""
    image_paths = sorted([os.path.join(folder_path, img) for img in os.listdir(folder_path) if img.endswith((".jpg", ".png"))])

    predictions_list = []

    for img_path in image_paths:
        # Load image and preprocess
        img = image.load_img(img_path, target_size=CONFIG["image_size"])
        img_array = image.img_to_array(img) / 255.0 # Normalize
        img_array = np.expand_dims(img_array, axis=0) # Add batch dimension

        # Get prediction
        prediction = model.predict(img_array)
        label = "AI-Generated" if prediction[0][0] > 0.5 else "Real"
        predictions_list.append((os.path.basename(img_path), label))

    # Save predictions to a CSV file
    df = pd.DataFrame(predictions_list, columns=["Image Name", "Predicted Label"])
    df.to_csv(output_csv, index=False)
    print(f"\nPredictions saved to {output_csv}")

    return predictions_list

# Run predictions on Test Dataset 1 and save results
predictions_test1 = predict_images_from_folder(testdataset1_path)

# Display results
print("\nPredicted Labels for Test Dataset 1:\n")
for img_name, label in predictions_test1:
    print(f"{img_name}: {label}")

# Optional: Plot images with predictions
plt.figure(figsize=(12, 6))
for i, (img_name, label) in enumerate(predictions_test1[:20]): # Show first 20 images
    img_path = os.path.join(testdataset1_path, img_name)
    img = image.load_img(img_path, target_size=CONFIG["image_size"])
    plt.subplot(4, 5, i + 1)
    plt.imshow(img)
    plt.title(f"{img_name}\n{label}")
    plt.axis("off")

plt.tight_layout()
plt.show()
```